

Pr. Olivier Gruber

`olivier.gruber@univ-grenoble-alpes.fr`

Laboratoire d'Informatique de Grenoble

Université de Grenoble-Alpes

- About processes
 - Create child processes (fork,execv)
 - Waiting for a child process (wait)
- About files and pipes
 - Files are persistent seekable streams of bytes
 - Pipes are for inter-process communication, a stream of bytes or packets
- About shells
 - Command line interpreter
 - Internal versus external commands
 - Piping and redirecting commands

- Tree of processes
 - Each process is uniquely identified (pid)
 - Linux manages its processes as a tree
 - Each process has a parent and may have children
- Processes can be programmatically
 - Created (fork,execv) and killed (kill)
 - Use the command “man” to have the details
 - Do not kill the process with pid==1

Duplicate a Process

4

- Fork function

- Duplicate the process calling the function “fork”
- Look at the man for more infos

\$ man -man 2 fork

Debugging: we remind you that it is possible to **attach gdb to a running process**, if you know the pid of that process, via the gdb command:

(gdb) attach *pid-of-the-process-to-attach-to*

Do not forget to **load the symbol table** for the executable file running in that process

(gdb) symbol *path-to-executable-file*

Note: Eclipse does support GNU toolchain via the CDT plugin

```
#include <unistd.h>

int main(int argc, char** argv) {

    int pid = fork();

    switch(pid) {
        case -1: /* error */
            perror("fork failed: ");
            break;
        case 0: /* child code */
            ...
            break;
        default: /* parent code */
            ...
            break;
    }
    return;
}
```

- The function "fork"
 - Duplicate the process calling the function “fork”
- Wait function
 - Waits for a state change⁽¹⁾ in a child process
 - Wait for the identified child process

\$ man -man 2 waitpid

```
#include <unistd.h>

int main(int argc, char** argv) {
    int pid = fork();
    int status;
    switch(pid) {
        case -1: /* error */
            perror("fork: ");
            exit(-1);
        case 0: /* child code */
            printf("Child exiting...");
            break;
        default: /* parent code */
            if (-1==waitpid(pid,&status,0)) {
                perror("waitpid: ");
                exit(-1);
            }
            break;
    }
    exit(0);
}
```

(1): A state change is considered to be: the child terminated; the child was stopped by a signal; or the child was resumed by a signal.

- The function "fork"
 - Duplicate the process calling the function “fork”
- Wait function
 - Waits for a state change⁽¹⁾ in a child process
 - Wait for the identified child process to exit
 - Wait for any child process

\$ man -man 2 wait

Note: the wait call is equivalent to:

```
waitpid(-1,&status, flags);
```

which can be **non-blocking** with the flag WNOHANG that forces the call to return immediately if no child has exited.

```
#include <unistd.h>

int main(int argc, char** argv) {
    int pid = fork();
    int status;
    switch(pid) {
        case -1: /* error */
            perror("fork: ");
            exit(-1);
        case 0: /* child code */
            printf("Child exiting...");
            break;
        default: /* parent code */
            pid = wait(&status);
            if (-1==pid) {
                perror("wait: ");
                exit(-1);
            }
            break;
    }
    exit(0);
}
```

(1): A state change is considered to be: the child terminated; the child was stopped by a signal; or the child was resumed by a signal.

Loading an executable file

7

- The function "execve"

- Loads and executes, in the process calling the function "execve", the program pointed to by filename

\$ man -man 2 execve

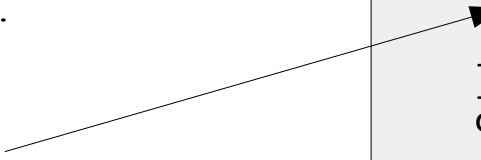
Note the real signature for main:

```
int main(int argc, char *argv[], char *envp[])
```

where both argv and envp are null-terminated arrays

Thus (argv[argc]==NULL) is true.

What does this call do,
with those arguments?



```
#include <unistd.h>

int main(int argc, char** argv,
        char *envp[]) {

    int pid = fork();
    int status;
    switch(pid) {
        case -1: /* error */
            perror("fork: ");
            exit(-1);
        case 0: { /* child code */
            char * args[3];
            args[0] = "ls";
            args[1] = "-al";
            args[2] = argv[1];
            execve("/bin/ls", args, envp);
            break;
        }
        default: /* parent code */
            if (-1==waitpid(pid,&status,0))
                perror("waitpid: ");
            break;
    }
    exit(0);
}
```

- **Reminder: file concept**
 - A file is a persistent sequence of bytes
 - Used a stream of bytes
- **Reminder: stream-related functions**
 - For more information, use "man"

\$ man -man 2 open

\$ man -man 2 close

\$ man -man 2 read

\$ man -man 2 write

```
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
```

```
int open(const char *pathname, int flags);
```

Flags: O_RDONLY, O_WRONLY, or O_RDWR.
These flags respectively mean read-only,
write-only, or read/write.

Use the flag O_CREAT to create a file
Use the flag O_TRUNC to truncate it

```
int close(int desc);
```

```
ssize_t read(int desc, void *ptr, size_t nb_octet);
```

```
ssize_t write(int desc, void *ptr, size_t nb_octet);
```


- A pipe is a inter-process communication mechanism
 - A pipe is a unidirectional stream of bytes between two processes
 - One process writes bytes into the pipe
 - The other process reads bytes from the pipe
 - For more information, use "man"

\$ man -man 2 pipe

```
#include <unistd.h>

int pipe(int pipefd[2]);

/*
 * The array pipefd is used to return two file
 * descriptors referring to the ends of the pipe.
 * pipefd[0] refers to the read end of the pipe.
 * pipefd[1] refers to the write end of the pipe.
 *
 * On success, zero is returned.
 * On error, -1 is returned, and errno is set
 * appropriately.
 */
```

About pipes and the shell

10

- Well, the shell provide a way to pipe commands
 - Via the character '|', like in these examples
- Requires a pipe in order to:
 - Connect the standard output stream of the command on the left
 - To the standard input stream of the command on the right
- Using the function "dup2"
 - The dup2() system call creates a copy of the file descriptor oldfd to the file descriptor newfd

\$ ls | grep toto

\$ cat toto | wc

```
#include <unistd.h>

int dup2(int oldfd, int newfd);
```

- Example using the function "dup2"

So, what does this program do?

```
#include <unistd.h>

int pipe[2];
#define READ_END 0
#define WRITE_END 1

int main(int argc, char** argv) {

    pipe(pipe);

    int pid = fork();

    switch(pid) {
        case -1: /* error */
            perror("panic: ");
            break;
        case 0: /* child code */
            dup2(pipe[READ_END],STDIN_FILENO);
            ...
            break;
        default: /* parent code */
            dup2(pipe[WRITE_END],STDOUT_FILENO);
            ...
            break;
    }
    return;
}
```

- Terminal window
 - The graphical user interface for a shell... running as a child process...
- A shell is just another program
 - It can be written in C, the needed functions are in the standard C library
 - It may print stuff through its standard output
 - It reads its standard input for the command lines typed by the end user
 - It interprets the command lines and executes their semantics
- Internal versus external commands
 - Some commands are internal commands: known to the shell implementation
 - Example: the command "cd" ou "pwd" are internal commands
 - Other commands are executable files on a "PATH"
 - Example: the command "ls" is in fact the program /bin/ls
 - The PATH is an environment variable, the environment variables are maintained by the shell
 - Example: PATH="/bin:/sbin:/home/droopy"